

Entrega 5: entrenamiento de redes neuronales

Aprendizaje Automatico y Big Data- Alejandro Barrachina Argudo

Introducción

En este documento se explicará el código del entregable 5 y el proceso de entrenamiento de redes neuronales.

Para esta práctica se usarán los siguientes *imports* vistos en la figura 0.1. Parte del código se reutiliza de la práctica anterior.

```
1 import numpy as np
2 import os
3 import scipy.io as sio
4 import scipy.optimize as opt
5 import utils
6 import matplotlib.pyplot as plt
7 from logistic_reg import sigmoid, plot_folder, csv_folder
```

Figura 0.1: Código de las bibliotecas usadas



Figura 0.2: Ejemplo de los dígitos del *dataset*

1. Entrenamiento de redes neuronales

Para comprobar la red neuronal utilizaremos la función *neural_network* implementada en la figura 1.1. Esta función se implementa para un número indeterminado de capas.

Reimplementamos la función de coste (figura 1.2) para que se adapte a las dos capas del ejercicio, esta función incluye también la regularización. La función de coste puede dar error para los números 0 y 1 por el uso del logaritmo, usamos la función *fix_data* (figura 1.3) para añadirle un infinitesimal valor para que no sea exactamente 0 o 1. Para la propagación primero haremos uso de la función de red neuronal para hacer la propagación hacia adelante y luego la propagación hacia atrás. Todo esto ocurre en la función *backprop* que se muestra en la figura 1.4.

Hacemos el descenso de gradiente en la función *gradient_descent* implementada en la figura 1.5. En esta función se hace uso de la función *backprop* para calcular los gradientes y actualizar los pesos.

Para terminar usamos *prediction*, *predict_percentage* y *random_init* como utilidades para el proceso de entrenamiento (figuras 1.6, 1.7 y 1.8).

```
1
2 def neural_network(X: np.ndarray, thetas: np.ndarray) -> tuple[np.ndarray, np.ndarray]:
3     """Generate the neural network with a given set of weights
4
5     Args:
6         X (np.ndarray): data
7         thetas (np.ndarray): array containing the weights for each layer
8
9     Returns:
10         tuple[np.ndarray, np.ndarray]: tuple containing the activations and the z values for
11         each layer
12     """
13     a = []
14     z = []
15     a.append(X.copy())
16     for theta in thetas:
17         a[-1] = np.hstack((np.ones((a[-1].shape[0], 1)), a[-1]))
18         z.append(np.dot(a[-1], theta.T))
19         a.append(sigmoid(z[-1]))
```

Figura 1.1: Función *neural_network*

```

1
2 def cost(theta1: np.ndarray, theta2: np.ndarray, X: np.ndarray, y: np.ndarray, lambda_: float
  = 0.0) -> float:
3     """
4     Compute cost for 2-layer neural network.
5
6     Parameters
7     -----
8     theta1 : array_like
9         Weights for the first layer in the neural network.
10        It has shape (2nd hidden layer size x input size + 1)
11
12    theta2: array_like
13        Weights for the second layer in the neural network.
14        It has shape (output layer size x 2nd hidden layer size + 1)
15
16    X : array_like
17        The inputs having shape (number of examples x number of dimensions).
18
19    y : array_like
20        1-hot encoding of labels for the input, having shape
21        (number of examples x number of labels).
22
23    lambda_ : float
24        The regularization parameter.
25
26    Returns
27    -----
28    J : float
29        The computed value for the cost function.
30
31    """
32    L = 2
33    layers = [theta1, theta2]
34    k: int = y.shape[1]
35    h, z = neural_network(X, [theta1, theta2])
36
37    h = h[-1]
38
39    h = fix_data(h)
40
41    J = y * np.log(h)
42    J += (1 - y) * np.log(1 - h)
43
44    J = -1 / X.shape[0] * np.sum(J)
45
46    if lambda_ != 0:
47        reg = 0
48        for layer in layers:
49            reg += np.sum(layer[:, 1:] ** 2)
50    J += lambda_ / (2 * X.shape[0]) * reg

```

Figura 1.2: Función de coste

```

1
2 def fix_data(X: np.ndarray) -> np.ndarray:
3     """Fixes the data to avoid log(0) errors
4
5     Args:
6         X (np.ndarray): train data
7
8     Returns:
9         np.ndarray: matrix with no 0 or 1 values
10    """

```

Figura 1.3: Función *fix_data*

```

1
2 def backprop(theta1: np.ndarray, theta2: np.ndarray, X: np.ndarray, y: np.ndarray, lambda_:
3     float) -> tuple[float, np.ndarray, np.ndarray]:
4     """
5     Compute cost and gradient for 2-layer neural network.
6
7     Parameters
8     -----
9     theta1 : array_like
10         Weights for the first layer in the neural network.
11         It has shape (2nd hidden layer size x input size + 1)
12
13     theta2: array_like
14         Weights for the second layer in the neural network.
15         It has shape (output layer size x 2nd hidden layer size + 1)
16
17     X : array_like
18         The inputs having shape (number of examples x number of dimensions).
19
20     y : array_like
21         1-hot encoding of labels for the input, having shape
22         (number of examples x number of labels).
23
24     lambda_ : float
25         The regularization parameter.
26
27     Returns
28     -----
29     J : float
30         The computed value for the cost function.
31
32     grad1 : array_like
33         Gradient of the cost function with respect to weights
34         for the first layer in the neural network, theta1.
35         It has shape (2nd hidden layer size x input size + 1)
36
37     grad2 : array_like
38         Gradient of the cost function with respect to weights
39         for the second layer in the neural network, theta2.
40         It has shape (output layer size x 2nd hidden layer size + 1)
41
42     """
43     m = X.shape[0]
44     L = 2
45
46     delta = np.empty(2, dtype=object)
47     delta[0] = np.zeros(theta1.shape)
48     delta[1] = np.zeros(theta2.shape)
49
50     a, z = neural_network(X, [theta1, theta2])
51
52     for k in range(m):
53         a1k = a[0][k, :]
54         a2k = a[1][k, :]
55         hk = a[2][k, :]
56         yk = y[k, :]
57
58         d3k = hk - yk
59         d2k = np.dot(theta2.T, d3k) * a2k * (1 - a2k)
60
61         delta[0] = delta[0] + \
62             np.matmul(d2k[1:, np.newaxis], a1k[np.newaxis, :])
63         delta[1] = delta[1] + np.matmul(d3k[:, np.newaxis], a2k[np.newaxis, :])
64
65     grad1 = delta[0] / m
66     grad2 = delta[1] / m
67
68     if lambda_ != 0:
69         grad1[:, 1:] += lambda_ / m * theta1[:, 1:]
70         grad2[:, 1:] += lambda_ / m * theta2[:, 1:]
71
72     J = cost(theta1, theta2, X, y, lambda_)

```

Figura 1.4: Función *backprop*

```

1
2 def gradient_descent(X: np.ndarray, y: np.ndarray, theta1: np.ndarray, theta2: np.ndarray,
3 alpha: float, lambda_: float, num_iters: int) -> tuple[np.ndarray, np.ndarray, np.ndarray]:
4     """Generates the gradient descent for the neural network
5
6     Args:
7         X (np.ndarray): Train data
8         y (np.ndarray): Expected output in one hot encoding
9         theta1 (np.ndarray): initial weights for the first layer
10        theta2 (np.ndarray): initial weights for the second layer
11        alpha (float): learning rate
12        lambda_ (float): regularization parameter
13        num_iters (int): number of iterations to run
14
15    Returns:
16        tuple[np.ndarray, np.ndarray, np.ndarray]: tuple with the final weights for the first
17        and second layer and the cost history
18    """
19    m = X.shape[0]
20    J_history = np.zeros(num_iters)
21    for i in range(num_iters):
22        print('Iteration: ', i + 1, '/', num_iters, end='\r')
23        J, grad1, grad2 = backprop(theta1, theta2, X, y, lambda_)
24        theta1 = theta1 - alpha * grad1
25        theta2 = theta2 - alpha * grad2
26        J_history[i] = J
27    print('Gradient descent finished.')
```

Figura 1.5: Función *gradient_descent*

```

1
2 def prediction(X: np.ndarray, theta1: np.ndarray, theta2: np.ndarray) -> np.ndarray:
3     """Generates the neural network prediction
4
5     Args:
6         X (np.ndarray): data
7         theta1 (np.ndarray): first layer weight
8         theta2 (np.ndarray): second layer weight
9
10    Returns:
11        np.ndarray: best prediction for each row in 'X'
12    """
13    m = X.shape[0]
14    p = np.zeros(m)
15    a, z = neural_network(X, [theta1, theta2])
16    h = a[-1]
```

Figura 1.6: Función *prediction*

```

1
2 def predict_percentage(X: np.ndarray, y: np.ndarray, theta1: np.ndarray, theta2: np.ndarray)
3 -> float:
4     """Gives the accuracy of the neural network
5
6     Args:
7         X (ndarray): Train data
8         y (ndarray): Expected output
9         theta1 (ndarray): First layer weights
10        theta2 (ndarray): Second layer weights
11
12    Returns:
13        float: Accuracy of the neural network
14    """
15    m = X.shape[0]
16    p = prediction(X, theta1, theta2)
```

Figura 1.7: Función *predict_percentage*

```

1
2 def random_init(size: tuple[int, int]) -> np.ndarray:
3     """Generates a random matrix of shape 'size' with values between -0.12 and 0.12
4
5     Args:
6         size (tuple[int,int]): shape of the generated matrix
7
8     Returns:
9         ndarray: random sample of shape 'size'
10    """

```

Figura 1.8: Función *random_init*

2. Flujo de entrenamiento

El programa llama a la función *test_learning* (figura 2.1) que se encarga de entrenar la red neuronal. Primero codificamos y a la codificación *one hot* mediante el uso de las funciones *oneHotEncoding* y *encoder* (figuras 2.2). Tras esto iniciamos los valores de iteración, número de capas, λ y α e iniciamos aleatoriamente θ_1 y θ_2 . Tras esto ejecutamos el descenso de gradiente, guardamos la matriz de confusión (usando la función de la práctica anterior) y guardamos los valores de la red neuronal usando la función *save_nn* (figura 2.3).

```

1
2 def test_learning(X: np.ndarray, y: np.ndarray) -> None:
3     """Tests the training of the neural network
4
5     Args:
6         X (np.ndarray): train data
7         y (np.ndarray): unencoded expected results
8     """
9     y_encoded = oneHotEncoding(y)
10    input_layer_size = X.shape[1]
11    hidden_layer_size = 25
12    num_labels = 10
13    lambda_ = 1
14    alpha = 1
15    num_iters = 1000
16
17    theta1 = random_init((hidden_layer_size, input_layer_size + 1))
18    theta2 = random_init((num_labels, hidden_layer_size + 1))
19
20    theta1, theta2, J_history = gradient_descent(
21        X, y_encoded, theta1, theta2, alpha, lambda_, num_iters)
22
23    save_nn(theta1, theta2, model_folder)
24
25    print(y)
26    print(prediction(X, theta1, theta2))
27
28    plot_confusion_matrix(y, prediction(
29        X, theta1, theta2), f'{plot_folder}confusion_matrix.png')
30
31    print('Expected accuracy: 95%. Got: ',

```

Figura 2.1: Función *test_learning*

```

1
2 def oneHotEncoding(y: np.ndarray) -> np.ndarray:
3     """Encodes the expected output to one hot encoding
4
5     Args:
6         y (np.ndarray): unencoded data
7
8     Returns:
9         np.ndarray: encoded data
10    """
11    def encoder(number: int) -> np.ndarray:
12        aux = np.zeros(10)
13        aux[number] = 1
14        return aux
15    y_encoded = [encoder(y[i] % 10) for i in range(y.shape[0])]

```

Figura 2.2: Función *oneHotEncoding*

```

1
2 def save_nn(theta1: np.ndarray, theta2: np.ndarray, folder: str) -> None:
3     """saves the neural network weights to a .mat file
4
5     Args:
6         theta1 (np.ndarray): first layer weights
7         theta2 (np.ndarray): second layer weights
8         folder (str): folder to save the mat
9     """

```

Figura 2.3: Función *save_nn*

3. Funciones auxiliares y resultados

Algunas funciones auxiliares del programa son las siguientes:

- *load_nn* (figura 3.2): Carga los valores de la red neuronal.
- *loadData* (figura 3.3): Carga los datos del *dataset*.
- *loadWeights* (figura 3.4): Carga los pesos de la red neuronal de ejemplo.
- *displayData* (figura 3.5): Muestra los dígitos del *dataset*.
- *main* (figura 3.6): Función principal del programa.

Podemos ver que la red neuronal consigue un 95% de acierto, representado en la figura 3.1.

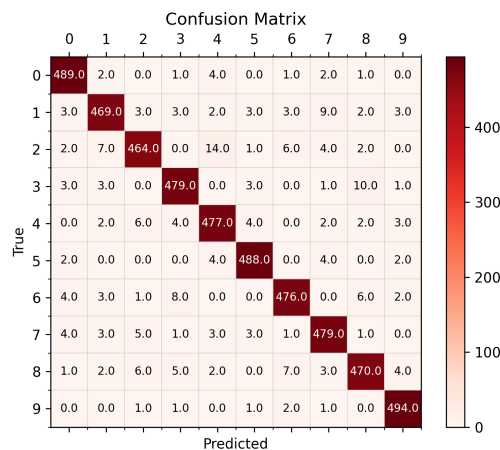


Figura 3.1: Resultado de la red neuronal

```

1
2 def load_nn(folder: str) -> tuple[np.ndarray, np.ndarray]:
3     """loads the neural network weights from a .mat file
4
5     Args:
6         folder (str): folder where the folder is stored
7
8     Returns:
9         tuple[np.ndarray, np.ndarray]: weights for the first and second layer
10    """
11    data = sio.loadmat(folder + 'nn.mat', squeeze_me=True)
12    theta1 = data['theta1']
13    theta2 = data['theta2']

```

Figura 3.2: Función *load_nn*

```

1
2 def loadData() -> tuple[np.ndarray, np.ndarray]:
3     """Loads the data from the .mat file
4
5     Returns:
6         tuple[np.ndarray]: X and y data
7     """
8     data = sio.loadmat('data/ex3data1.mat', squeeze_me=True)
9     X = data['X']
10    y = data['y']

```

Figura 3.3: Función *loadData*

```

1
2 def loadWeights() -> tuple[np.ndarray, np.ndarray]:
3     """Loads the example weights from the .mat file
4
5     Returns:
6         tuple[np.ndarray, np.ndarray]: example weights for the first and second layer
7     """
8     weights = sio.loadmat('data/ex3weights.mat', squeeze_me=True)
9     theta1 = weights['Theta1']
10    theta2 = weights['Theta2']

```

Figura 3.4: Función *loadWeights*

```

1
2 def displayData(X: np.ndarray) -> None:
3     """Displays the data in a 10x10 grid
4
5     Args:
6         X (np.ndarray): data to get a sample
7     """
8     if os.path.exists(f'{plot_folder}dataset.png'):
9         return
10    rand_indices = np.random.choice(X.shape[0], 100, replace=False)
11    utils.displayData(X[rand_indices, :])

```

Figura 3.5: Función *displayData*


```
1
2 def backprop_adapter(theta: np.ndarray, X: np.ndarray, y: np.ndarray, lambda_: float) -> float
3 :
4     theta1 = theta[:25 * 401].reshape(25, 401)
5     print(theta1.shape)
6     theta2 = theta[25 * 401:].reshape(10, 26)
7     print(theta2.shape)
8     # print(theta)
9     cost = backprop(theta1, theta2, X, y, lambda_)[0]
10    print(cost)
11    return cost
12
13 def minimize_test(X, y) -> None:
14     y_encoded = oneHotEncoding(y)
15     input_layer_size = X.shape[1]
16     hidden_layer_size = 25
17     num_labels = 10
18     lambda_ = 0.1
19     alpha = 1
20     num_iters = 1
21
22     theta1 = random_init((hidden_layer_size, input_layer_size + 1))
23     theta2 = random_init((num_labels, hidden_layer_size + 1))
24
25     pesos = np.random.rand(401 * 25 + 26 * 10) * 0.12 * -0.12
```

Figura 3.6: Función *main*